Environment

- Cluster with 3 nodes Nivdia T4 GPU
- Standard GKE cluster :region us-central1
- 10 GB ImageNet dataset dir:
- ResNet 50 pre-trained model

Cluster Creation

Used the following gcloud command to create a cluster named 'my-dask':

`gcloud container clusters create dask-cluster --num-nodes=3 --zone=us-central1 --disk-type=pd-standard --disk-size=10`

```
gongyitong@10-16-223-60 ~ % helm repo add dask https://helm.dask.org/
helm repo update

"dask" has been added to your repositories
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "dask" chart repository
Update Complete. ⎈Happy Helming!⎈
gongyitong@10-16-223-60 ~ %
```

Helm Installation of Dask

Installed Dask using Helm with the following command:

```
helm install my-dask dask/dask \
    --set scheduler.replicas=1 \
    --set worker.replicas=2
```

Check Cluster Status

Checked the cluster status using:

`kubectl get pods`

```
gongyitong@gongyitongMacBook-Pro data % kubectl get pods
NAME                              READY   STATUS    RESTARTS   AGE
my-dask-jupyter-558fcb5d46-6zwq4   1/1     Running   0          7h15m
my-dask-scheduler-5f778c9c4-4ltdk   1/1     Running   0          7h15m
my-dask-worker-64dd4775df-dtptb     1/1     Running   0          7h15m
my-dask-worker-64dd4775df-m5l8r     1/1     Running   0          7h15m
gongyitong@gongyitongMacBook-Pro data %
```

To learn more about the release, try:

Monitoring Cluster Status via External Interface

Accessed the cluster status through:

http://10.244.0.6:8786/status

Scheduler tcp://10.244.0.6:8786

Logs | Exceptions | Bokeh

Workers

Worker	Name	Cores	Memory	Memory Use	Occupancy	Processing	In-Memory	Services	Logs	Last Seen
tcp://10.244.0.3:36841	tcp://10.244.0.3:36841	4	1.94 GiB		0.00 us	0	0	dashboard	logs	134.23 ms
tcp://10.244.0.5:33911	tcp://10.244.0.5:33911	4	1.94 GiB		0.00 us	0	0	dashboard	logs	128.56 ms

Create Task Environment and Run Image

Built the Docker image with the following command:docker build -t image-processing

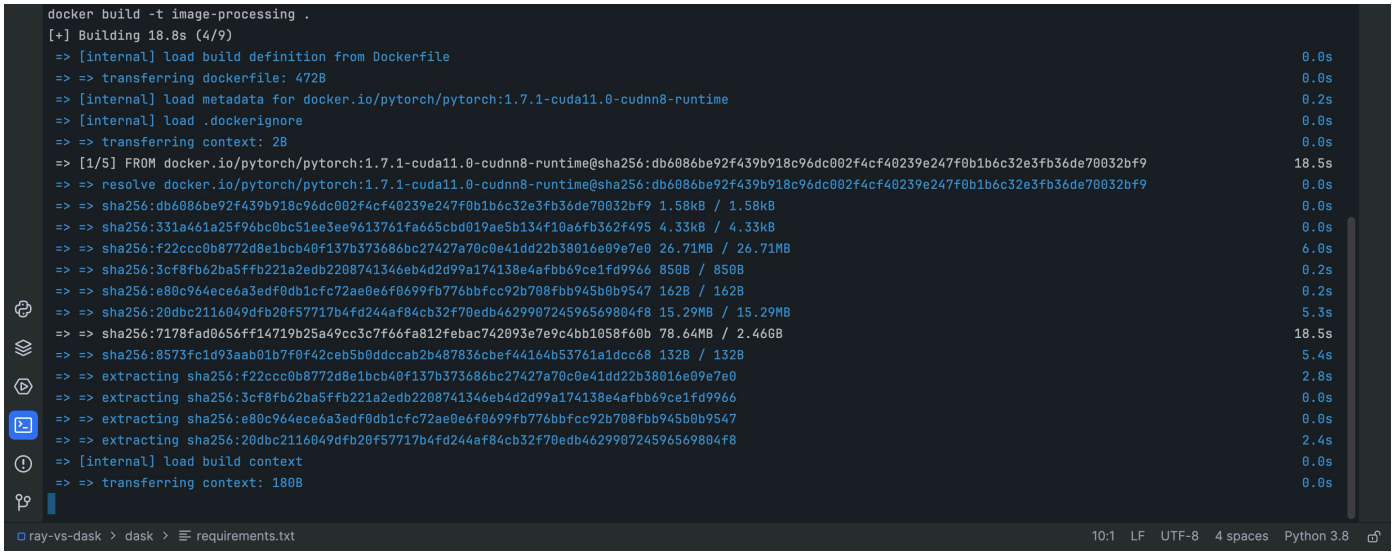

Image name: image-processing

Applied changes to the environment image to: `my-dask-scheduler.yaml` and `my-dask-worker.yaml`

Run Script in Environment

Connected to the scheduler using the Dask client with the following line of code:

```
client = Client('10.244.0.6:8786')  # Dask scheduler address and port
```

```
gongyitong@gongyitongMacBook-Pro dask % python3 main.py
Image processing completed in 987.00 seconds.
Average time per batch: 0.09885 seconds
Image processing results obtained.
```